

# Crawlle – A general purpose web crawler

## Features

- Custom priority/heuristic functionality to guide the crawler through the web
- Never go the same place twice!
- Massive parallelization, decide for yourself how far you want to go!
- Tree traversal back to seed from any discovered page
- Automatically store your results to a neat database (SQLite as of now)
- Download discovered files based on suffix, URL or regex matching
- Simple graphical user interface
- Automatically save your current session and all working data and resume later!
- Loosely coupled and easily maintainable architecture

## How to use

### Jar file

Download the jar, configure the properties file and run.

### Source

Download the source, the required libraries and run the Crawlle.java class.

### Required libraries

- guava-16.0.1
- jsoup-1.7.3
- sqlite-jdbc-3.7.2

## How it works

The crawler starts from a seed, then checks all the links on that page, and continues to the url that yielded the highest priority given the configuration. This can be searching for pdf-files, url that include “hello” etc. It will periodically store all the analyzed pages to a database. Due to the database becoming insanely heavy I’ve chosen not to include the source files. The crawler can also mark files/urls for download. For instance, one scenario could be: Go to hello.com and download all the images on that domain.

## Potential weaknesses

- Requires massive amounts of computer power if you intend to use it extensively
- High, monotonically increasing, memory usage due to having to cache discovered URLs in order to not enter a crawler loop

## Future improvements

The improvements and potential is highly dependent on the line of work the crawler is used for. However, some specifics might be:

- Implement a GUI for the configuration
- Use big data analytics to make sense of data collected
- Implement a relevance filter to prioritize pages based on content rather than URL

## How to expand

Expansion is easy. It’s already heavily parallelized with SWING (the GUI) being the biggest bottleneck. If instead of writing to a local database, decentralized instances could write to one common one. Depending on the implementation this could be very suitable for clusters and other multi- and/or super-computers.